

Statistics 243: *Assignment 3*

William J. De Meo

November 30, 1997

1 Cholesky Decomposition

Theorem:¹ A matrix A is symmetric positive definite (spd) if and only if there is a unique lower triangular nonsingular matrix L , with positive diagonal entries, such that $A = LL^t$.

$A = LL^t$ is called the *Cholesky decomposition* of A , and L is called the Cholesky factor of A . A simple algorithm for computing the Cholesky factor of an arbitrary spd matrix is found in the subroutine `cholesky()` in the Appendix. If the subroutine breaks down, the matrix passed to it was not spd.

2 Simulation

We construct a program for simulating groups of uncorrelated, as well as correlated data, and consider the test statistic for equality of means. It is observed that, when the data are correlated, this test statistic is far from the theoretical values assumed by an F-distribution.

An attempt was made to use the `unif()` and `normal()` subroutines from assignment 2 to generate normal (0,1) random numbers. Although these routines worked well for assignment 2, their use was abandoned after obtaining bad data values. A NAG routine was used instead to produce pseudo iid normal (0,1) random numbers. We now describe the technique used to generate correlated random numbers from the uncorrelated ones.

Suppose X has mean θ variance Σ . Then consider LX , which has

$$E(LX) = L\theta \quad \text{and} \quad V(LX) = L\Sigma L^t$$

Similarly, suppose X has mean 0 and variance I . Then $\tilde{X} = LX$ has variance $\Sigma = LL^t$. So we can create correlated random variables with variance covariance structure Σ by starting with uncorrelated random variables in X , taking the Cholesky decomposition $\Sigma = LL^t$, and then applying L to X .

Section 5.2 of the appendix lists the programs which carry out the foregoing. There is also listed the output from a sample run.

3 Gram-Schmidt orthogonalization (and other QR decompositions)

Our task is to write and test a function for the Gram-Schmidt orthogonalization (GS) of an arbitrary $n \times p$ matrix. Being arbitrary, we cannot assume that the matrix is *well-conditioned*², and it might seem desirable

¹A proof can be found in [1]

²We call a problem well-conditioned (ill-conditioned) if it has a small (large) condition number. In the present context, the condition number of the matrix X is $\|X\|_2\|X^{-1}\|_2$, and in general, the condition number can be defined heuristically as the inverse of the “distance” to the nearest singular problem. Consult, e.g., [1] for further details.

for our program to produce a warning when the data are *ill-conditioned*. However, we hold that this is unnecessary for the following reasons: In practice, all matrices are of full rank due to round-off error, and there exist algorithms that guarantee a numerically stable orthogonalization. Therefore, given any matrix, a good QR decomposition routine should proceed to numerically stable results. On the other hand, one might argue that, when using such algorithms to compute parameter estimates for least squares, as we do in the following problem, suppressing warnings of ill-conditioned data is misleading. This is a weak argument for a number of reasons. If the user of least squares estimates believes that the columns of X (covariates in the statistical setting) might be nearly linearly dependent, he should consider an SVD algorithm instead of GS. If he is too poor for SVD, then perhaps a rank revealing QR decomposition would be within the computing budget. The QR algorithm of this paper is *backward stable*³ and, using column pivoting, it is in some sense rank revealing in that small values on the diagonal of the resulting R matrix indicate the degree of ill-conditioning.

Any QR decomposition that uses a Gram-Schmidt type algorithm is unstable and, when the data are ill-conditioned, the resulting Q matrix may be far from orthogonal. Two QR decomposition procedures guarantee backward stability; one involves *Givens rotations*, and the other, *Householder reflections*. In this paper, we use the latter. Briefly, a Householder reflection (or transformation) $P = I - 2uu^t$ applied to x , is a reflection of x through the plane perpendicular to u . It is easy to verify that P is a symmetric orthogonal matrix. Why is this useful for QR? Given any vector x , we can find an Householder transformation P such that Px is a multiple of e_1 , where (e_1, e_2, \dots) is the standard basis. Therefore, we orthogonalize the columns of any matrix X by applying a sequence of Householder transformations P_i , where P_i annihilates elements $x(i+1, i), x(i+2, i), \dots$ of X . This results in an upper triangular⁴ matrix

$$\tilde{R} \equiv P_p \cdots P_1 X$$

Suppose $n > p$, and let Q_1 be the first p columns of the $n \times n$ orthogonal matrix $P_1 \cdots P_p \equiv (Q_1 \ Q_2)$, and R the first p rows of the $n \times p$ matrix

$$\tilde{R} \equiv \begin{pmatrix} R \\ O \end{pmatrix}$$

Then,

$$X = P_1 \cdots P_p \tilde{R} = Q_1 R \tag{1}$$

is a QR decomposition of X .

We could take each column of X in order, and perform the required Householder reflection, but pivoting can be used to make the algorithm more stable. Pivoting also has the nice side effect of producing an automatic estimate of the smallest singular value.⁵ Pivoting is introduced by simply selecting, of the remaining columns of X , the one with the largest norm as the next one to be reflected. To keep track of which vectors are where, we note that the column pivot is equivalent to right multiplication of X by a permutation of the identity. Since a permutation of the identity has its own transpose for an inverse, transforming back to the original ordering is easy.

To construct the Q and R matrices explicitly (Q is rarely required), we only need the Householder reflector vectors, u , and the upper triangle (trapezoid) of \tilde{R} . This is all that is returned (via the argument list) by the routine `qr`. The routine `qrpivot` requires the user to pass an additional matrix, E , which will equal the permutation matrix on exit (since it is essential that we know what permutation of our matrix was really decomposed).

The core of the algorithm just described is performed by the functions `qr()`, `qrpivot()`, and `House()`, each of which appears in the appendix. In their current state, these functions use only Level 2 BLAS.

³Informally, we say an algorithm is *backward stable* if it produces the exact answer to a slightly wrong problem.

⁴or, more accurately, upper trapezoidal when $n \neq p$

⁵The smallest singular value provides useful information about ill-conditioning. It is the norm of the smallest perturbation that can lower the rank of X .

However, it is possible to restructure the QR algorithm to take advantage of the Level 3 BLAS (as is done in the LAPACK routine `sgeqrf`).

4 Regression Using Householder Reflections

Given a vector Y (the response variable) and a matrix $X = [\mathbf{1} \ x_1 \ \dots \ x_p]$ (of intercept and covariates), we often make the following simplifying assumptions:

1. $Y = X\beta + \epsilon$
2. $\epsilon \sim N(\mathbf{0}, \sigma^2 \mathbf{I})$

We estimate the coefficients β with the usual OLS estimate given by the normal equations:

$$\hat{\beta} = (X^t X)^- X^t Y \quad (2)$$

where $(X^t X)^-$ is a generalized inverse of $X^t X$. The matrix $X^t X$ has the same rank as X , so if the our covariates are not *exactly* linearly dependent, then $(X^t X)^- = (X^t X)^{-1}$.

If the data are known to be well-conditioned, the fastest method of solving for $\hat{\beta}$ proceeds by simply taking the Cholesky decomposition and solving the normal equations (equations (2) above). Briefly, we would proceed as follows:

1. Do a Cholesky decomposition of $X^t X = LL^t$.
2. Let $\theta = L^t \hat{\beta}$. So that $L\theta = X^t y$.
3. Solve the l.t. system for θ .
4. Solve the l.t. system $L\hat{\beta} = \theta$ for $\hat{\beta}$.

However, that method is not very stable and can fail when the covariate matrix X is nearly rank deficient. Therefore, we will use the more stable QR algorithm with pivoting. From equation (1), we then have $XE = Q_1 R$. Therefore, $X^t X = E(R^t R)E^t$, and, using the fact that the permutation matrix E is orthogonal, (2) becomes

$$\hat{\beta} = ER^{-1}Q_1^t Y \quad (3)$$

The residuals are then

$$\begin{aligned} e &\equiv Y - X\hat{\beta} \\ &= Y - (Q_1 R E^t)(ER^{-1}Q_1^t Y) \\ &= (I - Q_1 Q_1^t)Y \end{aligned}$$

Recalling that $(Q_1 \ Q_2)$ is an $n \times n$ orthogonal matrix, we see that $(I - Q_1 Q_1^t) = Q_2 Q_2^t$. Whence,

$$e = Q_2 Q_2^t Y \quad (4)$$

Notice that (3) and (4) are obtained from transformations of Y , namely X 's Householder transformations

$$P_p \cdots P_1 Y = \begin{pmatrix} Q_1^t Y \\ Q_2^t Y \end{pmatrix}$$

Our routine `qrpivot` returns the arguments E and R , but not $(Q_1 \ Q_2)$. That is because construction of $(Q_1 \ Q_2)$ is unnecessary. Instead, recall that $P_j = I - 2u_j u_j^t$, so applying $(Q_1 \ Q_2)^t = P_p \cdots P_1$ to Y is

equivalent to the following algorithm:

```

for j = 1 to p
  a = -2 ∑i=jn uijYi
  for i = j to n
    Yi = Yi + auij
  end for
end for

```

For the regression problem, like many others, the u vectors are enough, and that is why the u 's, and not the Q 's, are returned by the QR routines.

The program **regress.c**, calls the functions **qrpivot()** and **reg()** to solve the least squares problem. The subroutine **reg** applies the orthogonal transformation $(Q_1 \ Q_2)^t$ to Y using the algorithm above, and applies ER^{-1} to the first p elements ($Q_1^t Y$) to find the solution $\hat{\beta}$. Then, again using the algorithm above but with j descending from p to 1, **reg** applies Q_2 to the last $n - p$ elements of $(Q_1 \ Q_2)^t Y$ by first annihilating its first p elements. More precisely,

$$\begin{aligned} P_1 \cdots P_p \begin{pmatrix} O \\ Q_2^t Y \end{pmatrix} &= (Q_1 \ Q_2) \begin{pmatrix} O \\ Q_2^t Y \end{pmatrix} \\ &= Q_2 Q_2^t Y \end{aligned}$$

By (4), this is the vector of residuals.

Finally, **regress.c** computes standard errors of the estimates. Note that equation (3) implies

$$\begin{aligned} \text{Var}(\hat{\beta}) &= \text{Var}(ER^{-1}Q_1^t Y) \\ &= (ER^{-1}Q_1^t) \text{Var}(Y) (ER^{-1}Q_1^t)^t \\ &= \sigma^2 ER^{-1} R^{-t} E^t \end{aligned}$$

which agrees with the usual formula under the i.i.d. normal assumption, $\sigma^2(X^t X)^{-1}$. So, to get standard error estimates, we only need the matrix ER^{-1} (which was derived when computing $\hat{\beta}$), and an estimate of σ^2 . To get an estimate of σ^2 we note that the sum of squared residuals is

$$\begin{aligned} \|Y - X\hat{\beta}\|^2 &= Y^t (I - X(X^t X)^{-1} X^t) Y \\ &= Y^t (I - Q_1 Q_1^t) Y \\ &= Y^t e \end{aligned}$$

So we take as our estimate of σ^2 the mean squared error, $Y^t e / (n - p)$.

A listing of the program **regress.c** and a test run using some wire bond strength data, appear in Appendix section 5.4.

5 Appendix

5.1 Cholesky Decomposition

A few lines of MATLAB can be used to produce a random spd matrix and compute the Cholesky factor:

```
>> n=5;
>> A=randn(n);
>> % Generate a random nxn spd matrix:
>> A = A'*A;
>> % Write the matrix A to datafile:
>> fid = fopen('datafile','w');
>>     fprintf(fid,'%f\n',A);
>>     fclose(fid);
>> % Perform Cholesky decomposition:
>> R = chol(A);
>> R'
```

ans =

1.5120	0	0	0	0
-0.1738	2.4651	0	0	0
-0.5637	0.5136	2.8050	0	0
-0.5603	-1.2035	-0.0434	2.1631	0
0.2969	1.7911	-0.6201	-1.0762	0.2292

We test out our program `cholesky.c` on the same matrix, now stored in `datafile`.

```
% cholesky
```

```
Enter file name containing the spd matrix: datafile
```

```
Enter its dimension: 5
```

```
The Cholesky factor is:
```

```
L =
1.51199
-0.17384      2.46512
-0.56367      0.51358      2.80495
-0.56033      -1.20354      -0.04342      2.16306
0.29689      1.79109      -0.62008      -1.07619      0.22917
```

The results are identical.

The listing for the main calling program `cholesky.c` is as follows:

```
/******
 * cholesky.c  main program for testing Cholesky      *
 * decomposition routine cholesky()                  *
 *                                                    *
 * Created by William J. De Meo                       *
 * on 11/29/97                                         *
```

```

*
*****/
#include <stdlib.h>
#include <stdio.h>
#include "prototypes.h"
#define MAX_NAME 100

void cholesky(long N, double *A, double *diag);
void read_name(char *);

main()
{
    char *filename;
    double *A, *diag;
    long i, j, dim;

    filename = calloc(MAX_NAME);

    printf("\nEnter file name containing the spd matrix: ");
    read_name(filename);
    printf("\nEnter its dimension: ");
    scanf("%d",&dim);

    A = dmalloc(dim*dim);
    diag = dmalloc(dim);

    matlabread(A, dim, dim, filename);
    /*matrix is stored contiguously column-wise */

    cholesky(dim,A,diag);

    printf("\nThe Cholesky factor is: \nL = \n");
    for(i=0;i<dim;i++)
    {
        for(j=0;j<i;j++)
            printf("%4.5lf \t", A[dim*j+i]);
        printf("%4.5lf", diag[i]);
        printf("\n");
    }
}

void read_name(char *name)
{
    int c, i = 0;

    while ((c = getchar()) != EOF && c != ' ' && c != '\n')
        name[i++] = c;
}

```

```

    name[i] = '\0';
}

```

The listing for the subroutine `cholesky()` is as follows:

```

/*~~~~~
cholesky.c

Created on 11/29/97 by William J. De Meo

Purpose: Cholesky decomposition of an n-by-n spd matrix
~~~~~*/
#include "prototypes.h"
#include <math.h>

/* Subroutine cholesky:

Arguments:
    N    dimension of A

    A
    on entry: the N by N matrix to be decomposed
    on exit: upper triangle is still A
            lower sub-triangle is the sub-triangle
            of the Cholesky factor L

    diag
    on entry: an arbitrary vector of length N
    on exit: the diagonal of the Cholesky factor L

*/

void cholesky(long N, double *A, double *diag)
{
    long i,j,k;
    for(j=0;j<N;j++)
        diag[j] = A[N*j+j];
    for(j=0;j<N;j++)
    {
        for(k=0;k<j;k++)
            diag[j] -= A[N*k+j]*A[N*k+j];
        diag[j] = sqrt(diag[j]);
        for(i=j+1;i<N;i++)
        {
            for(k=0;k<j;k++)
                A[N*j+i] -= A[N*k+i]*A[N*k+j];
            A[N*j+i]/=diag[j];
        }
    }
}

```

```
    }
}
```

5.2 Simulation

A trial run of the program `anova` produced the following:

```
% anova
```

```
How many groups? 5
```

```
How many in group 1? 6
```

```
How many in group 2? 6
```

```
How many in group 3? 6
```

```
How many in group 4? 6
```

```
How many in group 5? 6
```

```
How many simulated F's for this group structure? 10000
```

```
UNCORRELATED DATA
```

```
-----
```

```
90th percetile: theoretical = 2.184240, observed = 2.196409
```

```
95th percetile: theoretical = 2.758710, observed = 2.797849
```

```
99th percetile: theoretical = 4.177420, observed = 4.273534
```

```
CORRELATED DATA
```

```
-----
```

```
90th percetile: theoretical = 2.184240, observed = 14.493500
```

```
95th percetile: theoretical = 2.758710, observed = 18.310077
```

```
99th percetile: theoretical = 4.177420, observed = 29.474791
```

The output clearly indicates that when the data are correlated, the statistic no longer follows an F-distribution.

The program listing for `anova.c`, including the subroutines `F()` and `normal()` (even though I was forced to abandon `normal()`) appears below:

```
/*~~~~~
  anova.c
```


Created by William J. De Meo
on 11/30/97

Purpose: Simulating F-statistics for one way ANOVA
using correlated and uncorrelated data

```
~~~~~*/
#include <math.h>
#include <stdio.h>
#include "prototypes.h"

/* NAG prototypes */
double g01fdf_(double *p, double *DF1, double *DF2, int *ifail);
void m01caf_(double *RV, long *M1, long *M2, char *ORDER, int *ifail);
void g05fdf_(double *mean, double *sd, long *n, double *g);

/* BLAS prototypes */
/* C <- (alpha)AB + (beta)C */
void dgemm_(char *TRANSA, char *TRANSB, long *M, long *N, long *K, double *alpha,
            double *A, long *LDA, double *B, long *LDB, double *beta, double *C, long *LDC);

void F(long N, long *n, long k, double *AVE, double *ave, double *var, double *eff);

extern long I = (long)0;

main()
{
    char ORDER = 'A'; /* F stats will be sorted in ascending order */
    int ifail = 0;
    long k, i, j, p, q, r, N=(long)0, M = (long)12, numF, f, one = (long)1;
    double *eff, *pureff, *u, *g, *gtemp, *AVE, *VAR, *ave, *var, *sig, *diag, *work;
    double minimum=(double)1000, maximum=(double)-1000, P1, P2, P3, unit=(double)1, zero=(double)0;
    long *X, *n;

    /* BLAS arguments */
    double alpha = (double)1, beta = (double)0;
    char NOTRANS = 'N';

    AVE = dmalloc((long)1);
    VAR = dmalloc((long)1);
    gtemp = dmalloc((long)2);
    X = lmalloc((long)1);
    *X = time('\0');

    printf("\nHow many groups? ");
    scanf("%d", &k);
}
```

```

n = lmalloc(k); /* n vector stores number in each group */
ave = dmalloc(k);
var = dmalloc(k);

for(i=0;i<k;i++)
{
    printf("\nHow many in group %d? ", i+1);
    scanf("%d",&n[i]);
    N+=n[i];
}
M = 2*N; /* unifs fail about 27% of the time, so take twice as many */
g = dmalloc(N+1); /* normals will be stored in g (N+1 in case N odd)*/
u = dmalloc(M);

printf("\nHow many simulated F's for this group structure? ");
scanf("%d",&numF);
eff = dmalloc(3*numF); /* matrix of F-stats with df's */
pureff = dmalloc(numF); /* vector of F-stats (without df's) */

for(f=0;f<numF;f++)
{
    I = (long)0;
    for(r=0;r<M;r++)
        u[r] = unif(X);

    i=(long)0;
    p=(long)0;

    for(j=0;j<k;j++)
    {
        /*Forced to scrap the following broken random number generating code */
        /*~~~~~*/
        if(0)
        {
            while(i<n[j])
            {
                if(normal(u, M, gtemp)==1)
                {
                    g[p+i] = gtemp[0]; i++;
                    g[p+i] = gtemp[1]; i++;
                }
                else /* didn't get enough normals -- need new uniforms */
                {
                    printf("\n\nGenerating different unif(0,1) variables...\n\n");
                    I = 0;
                    for(r=0;r<M;r++)
                        u[r] = unif(X);
                }
            }
        }
    }
}

```

```

        }
    }
    if(i==n[j])/* i.e. even number in group */
        i=(long)0;
    else i=(long)1; /* i.e. odd number in group */
    /*then give last random number to the next group */
}
/*~~~~~*/

/* NAG routine puts normal(0,1) random numbers in g */
g05fdf_(&zero,&unit,n+j,g+p);

/* compute group averages and variances */
cmoment(g+p, n[j], ave+j, var+j);
/*printf("\nave[%d] = %lf, var[%d] = %lf\n\n",j+1,ave[j],j+1,var[j]);*/

p+=n[j];
}

cmoment(g,N,AVE,VAR); /* Get overall average */

/* finally compute the associated F-statistic */
F(N,n,k,AVE,ave,var,eff+3*f);
}

/* strip degrees of freedom from eff (so we can sort it) */
for(j=0;j<numF;j++)
    pureff[j] = eff[j*3];

/* sort the pureff vector of F statistics */
m01caf_(pureff, &one, &numF, &ORDER, &ifail);
printf("\n\nUNCORRELATED DATA\n");
printf("-----\n");
P1 = (double).9;
P1 = g01fdf_(&P1,eff+1,eff+2,&ifail);
printf("\n90th percetile: theoretical = %lf, observed = %lf\n",P1,pureff[(long)(.9*numF)]);
P2 = (double).95;
P2 = g01fdf_(&P2,eff+1,eff+2,&ifail);
printf("\n95th percetile: theoretical = %lf, observed = %lf\n",P2,pureff[(long)(.95*numF)]);
P3 = (double).99;
P3 = g01fdf_(&P3,eff+1,eff+2,&ifail);
printf("\n99th percetile: theoretical = %lf, observed = %lf\n",P3,pureff[(long)(.99*numF)]);

/*~~~~~
Correlated data
WARNING: This part only written for non-ragged arrays
i.e. all k group sizes must be equal*/

```

```

/* Test to see that all group sizes equal */
for(j=0;j<k;j++)
    if(n[0]!=n[j])
    {
        printf("\n\nGroups of different sizes\n");
        printf("\n Can't perform correlation simulations.\n");
        exit(0);
    }

/* now just refer to group size as *n */

sig = dmalloc(*n * *n);
diag = dmalloc(*n);
work = dmalloc(N);

for(f=0;f<numF;f++)
{
    /* NAG routine puts normal(0,1) random numbers in g */
    g05fdf_(&zzero,&unit,&N,g);

    for(j=0;j<*n;j++)
        for(i=j;i<*n;i++) /* only need lower triangle of symmetric matrix */
            sig[*n * j+i]=pow(0.7,(i-j));
    for(i=0;i<*n;i++)
        for(j=i+1;j<*n;j++)
            sig[*n * j +i]=0;

    cholesky(*n, sig, diag);
    for(j=0;j<*n;j++) sig[*n * j +j]=diag[j];
    free(diag);

    /*work <- sig*g */
    dgemm_(&NOTRANS, &NOTRANS, n, &k,n, &alpha, sig, n, g, n, &beta, work, n);
    for(j=0;j<N;j++) g[j] = work[j];

    /* compute group averages and variances */
    p=(long)0;
    for(j=0;j<k;j++)
    {
        cmoment(g+p, n[j], ave+j, var+j);
        p+=n[j];
    }
    /*printf("\nave[%d] = %lf, var[%d] = %lf\n\n",j+1,ave[j],j+1,var[j]);*/

    cmoment(g,N,AVE,VAR); /* Get overall average */

    /* finally compute the associated F-statistic */
    F(N,n,k,AVE,ave,var,eff+3*f);

```

```

    }

    /* strip degrees of freedom from eff (so we can sort it) */
    for(j=0;j<numF;j++)
        pureff[j] = eff[j*3];

    /* sort the pureff vector of F statistics */
    m01caf_(pureff, &one, &numF, &ORDER, &ifail);
    printf("\n\nCORRELATED DATA\n");
    printf("-----\n");
    printf("\n90th percetile: theoretical = %lf, observed = %lf\n",P1,pureff[(long)(.9*numF)]);
    printf("\n95th percetile: theoretical = %lf, observed = %lf\n",P2,pureff[(long)(.95*numF)]);
    printf("\n99th percetile: theoretical = %lf, observed = %lf\n",P3,pureff[(long)(.99*numF)]);
}

/* subroutine F() ~~~~~*/

Purpose:  Compute test statistic for equality of means

Arguments:

    N  total number of observations

    n  a vector of length k where the ith element
        contains the number of observations in the ith group

    k  the number of groups

    AVE  average of all observations

    ave  a vector of length k where the ith element
        contains the average of the ith group of observations

    var  a vector of length k where the ith element
        contains the empirical variance (mse) of the ith group of observations

    F  on entry: a vector of length 3
        on exit: first element is the F-statistic
                 second element is k-1, third element is N-k
~~~~~*/

void F(long N, long *n, long k, double *AVE, double *ave, double *var, double *F)
{
    long i,j;
    double den=(double)0,num=(double)0;

    F[1]=k-1; F[2]=N-k;

```

```

    for(i=0;i<k;i++)
    {
        num += n[i]*(ave[i]-*AVE)*(ave[i]-*AVE)/F[1];
        den += (n[i]-1)*var[i]/F[2];
    }

    F[0] = num/den;
}
/*~~~~~
normal.c

Subroutine for constructing normal(0,1) random numbers

Arguments:

    u  a vector of uniform random variables

    n  the length of u

    x  on entry:  an arbitrary vector of length at least 2
        on exit:  first two elements are two normal random variables

Return values:

    0  failure
        (not enough uniform random variables to construct 2 normals)

    1  success
~~~~~*/
#include <math.h>

int normal(double *u, long n, double *x)
{
    static int numcount = (long)0, dencount = (long)0;
    int count = (long)0;
    double s;
    extern long I;
    dencount++;

    do{

        /* The following lines display the proportion of times we are rejecting
           s, which we expect to be around .27
        if(count>0)
        {

```

```

        numcount++;
        printf("\n\n s = %lf, prop = %lf\n\n",
              s, (double)numcount/(double)dencount);
    }
    */
    if(I > n-2)
        return(0);/* set external I back to 0 in main() */
    s = (2*u[I] - 1)*(2*u[I] - 1) + (2*u[I+1] - 1)*(2*u[I+1] - 1);
    I += 2;
    count++;
}while(s >= (double)1);

x[0] = (2*u[I] - 1)*sqrt(-2*log(s)/s);
x[1] = (2*u[I+1] - 1)*sqrt(-2*log(s)/s);

return(1);

}

```

5.3 QR Decomposition

Using the MATLAB program `tests.m` (listed below), we produce a data file called `datafile` containing a random 4×3 matrix with condition number 100. The matrix is stored column-wise by MATLAB, so `datafile` contains the following:

```

-55.201723
-40.707641
42.683791
16.618428
-35.914110
-24.880338
28.457431
15.289218
20.047960
19.983140
-6.967239
75.587311

```

That is, the matrix to be decomposed is:

$$A = \begin{pmatrix} -55.201723 & -35.914110 & 20.047960 \\ -40.707641 & -24.880338 & 19.983140 \\ 42.683791 & 28.457431 & -6.967239 \\ 16.618428 & 15.289218 & 75.587311 \end{pmatrix}$$

We can use the program `QR` (which calls the subroutine `qr()`) to decompose A , as the following output demonstrates:

```
% QR
```

Enter file name containing the matrix: datafile

Enter the number of rows: 4

Enter the number of columns: 3

The orthogonalization produced:

82.47679	54.12546	-11.65654
-0.27012	-4.78269	-77.59750
0.28324	0.06872	20.14904
0.11027	0.55963	0.62361

with leading u's:

-0.91359	0.82589	-0.78174
----------	---------	----------

Running the program QRpiv (which calls the subroutine qrpivot()) on the same matrix produces:

% QRpiv

Enter file name containing the matrix: datafile

Enter the number of rows: 4

Enter the number of columns: 3

The orthogonalization produced:

82.47679	-11.65654	54.12546
-0.27012	-80.17079	-4.62918
0.28324	0.02372	-1.20202
0.11027	0.65823	-0.59810

With permutation matrix:

1.00000	0.00000	0.00000
0.00000	0.00000	1.00000
0.00000	1.00000	0.00000

and leading u's:

-0.91359	0.75244	0.80142
----------	---------	---------

We can check that these results are accurate using MATLAB (which also performs QR using Householder reflections). The QR portion of the MATLAB program `tests.m` produces the following:

```
>> % Perform QR without pivoting:
```

```
>> [Q,R] = qr(A);
```

```
>> R
```

```
R =
```

```
82.4768  54.1255 -11.6565
```



```

      0   -4.7827  -77.5975
      0         0   20.1490
      0         0         0

```

```

>> % Perform QR with pivoting:
>> [Qpiv, Rpiv, E] = qr(A);
>> Rpiv

```

```

Rpiv =

```

```

 82.4768  -11.6565   54.1255
      0  -80.1708  -4.6292
      0         0  -1.2020
      0         0         0

```

```

>> E

```

```

E =

```

```

 1     0     0
 0     0     1
 0     1     0

```

We see that the only difference in the results is that our subroutines are storing the u 's in the lower part of the R matrix.

The MATLAB program listing for `tests.m` is as follows:

```

% MATLAB code tests.m
% Created by William J. De Meo
% on 11/28/97
%
% Purpose: Perform QR decomposition (with and without pivoting)
%          on a random matrix of user specified dimension and
%          condition number
%
% Inputs:
%
%   m, n = numbers of rows, columns in test matrices
%          m should be at least n
%
%   cnd = condition number of test matrices to generate
%          (ratio of largest to smallest singular value)
%          cnd should be at least 1
%
% Outputs:
%
%   datafile = A file containing the matrix tested
%             (so that we can run QR and QRpiv on the same matrix)
%

```

```

% R = the R from the QR decomposition of A
%
% Rpiv, E = the R and permutation matrix E from QR decomposition
%           of A with pivoting
%
% Generate random matrix A, starting with the SVD of a random matrix
A=randn(m,n);
[u,s,v]=svd(A);
% Let singular values range from 1 to cnd, with
% uniformly distributed logarithms
sd = [1, cnd, exp(rand(1,n-2)*log(cnd))];
s = diag(sd);
A=u(:,1:n)*s*v';
%
% Write the matrix A to datafile:
fid = fopen('datafile','w');
    fprintf(fid,'%f\n',A);
    fclose(fid);

% Perform QR without pivoting:
[Q,R] = qr(A);
R
% Perform QR with pivoting:
[Qpiv, Rpiv, E] = qr(A);
Rpiv
E
%
% END tests.m

```

The following is the program listing for QRpiv.c:⁶

```

/*****
 * QRpiv.c  main program for testing QR
 * Orthogonalization routine qrpivot()
 *
 * Created by William J. De Meo
 * on 11/23/97
 *
 * Note: differences between QRpiv.c and QR.c are marked
 *       with the comment "QRpiv.c"
 *****/

#include <stdlib.h>
#include <stdio.h>
#include "prototypes.h"
#define MAX_NAME 100

```

⁶QR.c is not listed as it is almost identical to QRpiv.c. The only difference being that it calls `qr()` instead of `qrpivot()`. The differing lines are marked in QRpiv.c with the comment `/* QRpiv */`.

```

/* QRpiv.c */
void qrpivot(long M, long N, double *A, double *E, double *leadu);

void read_name(char *);

main()
{
    char *filename;
    double *x, *leadu, *E;
    long i, j, nrow, ncol, mindim;

    filename = calloc(MAX_NAME);

    printf("\nEnter file name containing the matrix: ");
    read_name(filename);
    printf("\nEnter the number of rows: ");
    scanf("%u",&nrow);
    printf("\nEnter the number of columns: ");
    scanf("%u",&ncol);
    mindim = lmin(nrow,ncol); /* mindim is the smaller dimension */

    x = dmalloc(nrow*ncol);
    leadu = dmalloc(mindim);
    E = dmalloc(ncol*ncol);

    matlabread(x, nrow, ncol, filename);
    /*matrix is stored contiguously column-wise */

    /* Test qrpivot: */
    qrpivot(nrow,ncol,x,E,leadu); /* QRpiv.c */

    printf("\nThe orthogonalization produced: \n");
    matprint(x,nrow,ncol);
    printf("\nWith permutation matrix: \n"); /* QRpiv.c */
    matprint(E,ncol,ncol); /* QRpiv.c */
    printf("\nand leading u's:\n");
    for(i=0;i<mindim;i++)
        printf("%4.5lf \t", leadu[i]);
    printf("\n");
}

void read_name(char *name)
{
    int c, i = 0;

    while ((c = getchar()) != EOF && c != ' ' && c != '\n')
        name[i++] = c;
    name[i] = '\0';
}

```

```
}

```

The functions `qr()`, `qrpivot()`, and `House()` are found in `House.c`, which is listed below:

```
/*~~~~~
House.c

Created on 11/12/97 by William J. De Meo
Last modified: 11/28/97

Purpose: QR decomposition of an m-by-n matrix using
Householder reflections

Further Details: This implementation uses BLAS 2
(matrix-vector mult. and rank 1 updates)

Dependencies: Requires subroutines found in the libraries:
sunperf, and blas
the later two are linked with the options:
-lsunperf -dalign -lblas
compilation must be done with the -dalign option
~~~~~*/
#include "prototypes.h"

/*~~~~~
BLAS Subroutine prototypes
~~~~~*/
double dnorm2_(long *N, double *x, long *INC); /* L2 norm of x */
double dcopy_(long *N, double *X, long *INCX, double *Y, long *INCY); /* y <- x */
double dgemv_(char *TRANSA, long *M, long *N, double *alpha, double *A, long *LDA,
double *x, long *INCX, double *beta, double *y, long *INCY);
/* y <- (alpha)Ax + (beta)y (or A^t if TRANSA='T') */

void dger_(long *M, long *N, double *alpha, double *x, long *INCX, double *y, long *INCY,
double *A, long *LDA); /* Rank 1 update A <- (alpha)xy^t + A */
/*~~~~~*/

/* Subroutine qr:

Arguments:
M (long) number of rows of A

N (long) number of columns of A

A (pointer to double)
on entry: the M by N matrix to be decomposed
on exit: upper-right-triangle = R
column i of lower trapezoid = a(i+1:M,i) = u(1:m,i)
where P(i) = I - 2 u(1:m,i)u^t(1:m,i) is the ith

```

Householder transformation (of dimension $(M-i) \times (M-i)$)

leadu (pointer to double)
 on entry: an arbitrary vector of length $\min(M-1, N)$
 on exit: the leading entries of the Householder vectors $u(i)$
 i.e. $u(i) = (\text{leadu}(i), a(i+1:M, i))$ $i=1, \dots, \min(M-1, N)$

Note that Q is obtained from augmenting the Householder transformations to be of proper dimensions, and then multiplying:

If $P'(i)$ denotes augmented $P(i)$,

$Q = P'(1) P'(2) P'(3) \dots =$

	1 0		1			
P(1)	-----		1	0		
	0 P(2)		-----		...	
			0	P(3)		

But this is left to the calling function and is not performed in `qr()`.

*/

```
void qr(long M, long N, double *A, double *leadu)
{
    char T = 'T';
    long i, j, nrow, ncol, mindim;
    double unit, zero, alpha=(double)-2;
    double *u, *y; /* used for temporary work space */
    long INC=1; /* INC is used to represent storage
                  spacing between elements */
    unit = (double)1; zero = (double)0;

    u = dmalloc(M); /* work space */
    y = dmalloc(N);

    mindim = lmin(M-1, N);

    for(i=0; i<mindim; i++)
    {
        nrow=M-i;
        ncol=N-i;

        House(nrow, A+(M*i+i), u);

        /* y <- (A^t)u (is working)*/
        dgemv_(&T, &nrow, &ncol, &unit, A+(M*i+i),
              &M, u, &INC, &zero, y, &INC);

        /* Rank 1 update: A <- A + (-2)uy^t i.e. A - 2uu^tA */
    }
}
```

```

    dger_(&nrow,&ncol, &alpha, u, &INC, y, &INC, A+(M*i)+i, &M);
    leadu[i] = u[0];

    /* store u(2:nrow) in A(i+1:M,i) */
    nrow--;
    dcopy_(&nrow, u+1, &INC, A+(M*i)+i+1,&INC);
}
    free(u); free(y);
}

/* Subroutine qrpivot:

Arguments: same as qr() with one exception:

        E (pointer to double)
            on entry: an arbitrary N by N matrix
            on exit: the permutation matrix
            The final decomposition is AE = QR
*/

void qrpivot(long M, long N, double *A, double *E, double *leadu)
{
    char T = 'T';
    long i,j, nrow, ncol, mindim, perm=0;
    double unit, zero, alpha=(double)-2, maxnorm, norm;
    double *u, *y; /* used for temporary work space */
    long INC=1;      /* INC is used to represent storage
                        spacing between elements */
    unit = (double)1; zero = (double)0;

    u = dmalloc(M); /* work space */
    y = dmalloc(N);
    /* Start permutation matrix as the identity */
    for(i=0;i<N;i++){
        E[N*i+i] = (double)1;
    }
    mindim = lmin(M-1,N);

    for(i=0;i<mindim;i++)
    {
        nrow=M-i;
        ncol=N-i;

        /* column pivot */
        maxnorm=0;
        for(j=i;j<N;j++)
        {

```

```

        norm = dnorm2_(&nrow,A+(M*j+i),&INC);
        if(norm>maxnorm)
        {
            perm=j;
            maxnorm=norm;
        }
    }

    if(perm>i)
    { /* If the i'th column was not the largest in norm, permute cols
       of A, and note it by swapping cols of pivot matrix*/
        dswap_(&M,A+(M*i),&INC,A+(M*perm),&INC);
        dswap_(&N,E+(N*i),&INC,E+(N*perm),&INC);
    }

    House(nrow, A+(M*i+i), u);

    /* y <- (A^t)u */
    dgemv_(&T, &nrow, &ncol, &unit, A+(M*i)+i,
           &M,u,&INC,&zero,y,&INC);

    /* Rank 1 update: A <- A + (-2)uy^t   i.e. A - 2uu^tA */
    dger_(&nrow,&ncol, &alpha, u, &INC, y, &INC, A+(M*i)+i, &M);
    leadu[i] = u[0];

    /* store u(2:nrow) in A(i+1:M,i) */
    nrow--;
    dcopy_(&nrow, u+1, &INC, A+(M*i)+i+1,&INC);
}
free(u); free(y);
}

void House(long N, double *a, double *u)
{
    double sign = (double)1, norm;
    long inc=1;
    long i;

    if(a[0] <= (double)0) sign = -1;

    /* u <- A(i:M,i) */
    dcopy_(&N, a, &inc, u, &inc);

    norm = dnorm2_(&N, a, &inc);          /* L2 norm of a = A(i:M,i) */
    u[0] += sign*norm;
    norm = dnorm2_(&N, u, &inc);          /* L2 norm of new u */

```

```

/* could also try the relation: norm = sqrt(2*(norma*norma + sign*u[1]*norma)) */

for(i=0;i<N;i++) u[i] /= norm;
/* consider skipping this normalization and putting the norms in the coefficient:
   alpha = 2/(unorm * unorm) */
}

```

5.4 Regression Using Householder Reflections

The file `datafile` contains 30 observations of 5 variables related to wire strength (plus a column of 1's for the intercept). The data were taken from a Stat 215a lab and can be found in `/saruman/class/data/s215/fall97/lab3.data`. Using the `regress` program, we can estimate the OLS coefficients and compute the statistics of interest as follows:

```
% regress
```

```
Enter file name containing the matrix [X,y]
(with intercept column included if desired): datafile
```

```
Enter the number of observations: 30
```

```
Enter the number of parameters (including intercept): 5
```

```
Rough estimate of smallest singular value of X:
R(5,5) = 0.393750
```

```
MSE = 26.604747
```

COEFFICIENT	SE
-37.47667	13.09964
0.21167	0.21057
0.49833	0.07019
0.12967	0.04211
0.25833	0.21057

OBS	RESIDUAL
1	0.86000
2	-1.15667
3	-0.49000
4	-2.70667
5	6.77667
6	1.56000
7	2.02667
8	-11.09000
9	-1.32333
10	-6.64000
11	-4.27333
12	7.11000

13	5.09333
14	-4.22333
15	-6.35667
16	4.52667
17	-1.09000
18	6.37667
19	-0.75667
20	5.34333
21	2.57667
22	-0.89000
23	-0.12333
24	0.81000
25	-1.90667
26	-0.80667
27	1.89333
28	7.59333
29	-10.60667
30	1.89333

The upper bound on the smallest singular value given by $R(5,5) = 0.393750$ indicates that X is at most a distance of 0.393750 from the nearest rank deficient matrix. Perhaps one of our variables is nearly constant and we don't need an intercept. The **regress** program handles this, requiring only that we modify the input file. The new data file, **newdata**, doesn't contain a column of 1's. The new results are as follows:

```
% regress
```

```
Enter file name containing the matrix [X,y]
(with intercept column included if desired): newdata
```

```
Enter the number of observations: 30
```

```
Enter the number of parameters (including intercept): 4
```

```
Rough estimate of smallest singular value of X:
```

```
R(4,4) = 25.868210
```

```
MSE = 33.956565
```

COEFFICIENT	SE
-0.12727	0.19667
0.41763	0.07261
0.05220	0.03644
0.06466	0.22527

OBS	RESIDUAL
1	-5.14401
2	-3.77132
3	-4.07304

```

4      -2.90034
5      4.64621
6      2.81891
7      2.31718
8      -7.41012
9      -5.39057
10     -7.31787
11     -5.91960
12     8.85310
13     4.89966
14     -1.02765
15     -4.12937
16     10.14332
17     -4.67304
18     9.57235
19     -3.37132
20     7.57063
21     -1.49057
22     2.78988
23     -2.25379
24     2.55310
25     -2.10034
26     -1.00034
27     1.69966
28     7.39966
29     -10.80034
30     1.69966

```

That looks better. A quick check can be made using one line of MATLAB:

```
>> [B,se] = lscov(X,Y,eye(30))
```

B =

```

-0.1273
 0.4176
 0.0522
 0.0647

```

se =

```

0.1967
0.0726
0.0364
0.2253

```

The results are identical (not a coincidence – MATLAB least squares functions also use Householder QR with pivoting).

The following is the program listing for `regress.c`:

```

/*****
 * regress.c  main program for performing regression      *
 *                                                    *
 * Created by William J. De Meo                        *
 * on 11/28/97                                          *
 *                                                    *
 *****/

#include <stdlib.h>
#include <stdio.h>
#include <math.h>
#include "prototypes.h"
#define MAX_NAME 100

/*~~~~~
  BLAS Subroutines prototypes */
void dcopy_(long *N, double *X, long *INCX, double *Y, long *INCY); /* y <- x */
double ddot_(long *N, double *X, long *INCX, double *Y, long *INCY); /* returns x * y */
void dtrsv_(char *UPLO, char *TRANSA, char *DIAG, long *N, double *A,
            long *LDA, double *Y, long *INCY); /* y <- inv(A)*y */

/* y <- (alpha)Ax + (beta)y (or A^t if TRANSA='T') */
void dgemv_(char *TRANSA, long *M, long *N, double *alpha, double *A, long *LDA,
            double *x, long *INCX, double *beta, double *y, long *INCY);

/* C <- (alpha)AB + (beta)C */
void dgemm_(char *TRANSA, char *TRANSB, long *M, long *N, long *K, double *alpha,
            double *A, long *LDA, double *B, long *LDB, double *beta, double *C, long *LDC);
/* B <- alpha*inv(A)*B */
void dtrsm_(char *SIDE, char *UPLO, char *TRANSA, char *DIAG, long *M, long *N,
            double *alpha, double *A, long *LDA, double *B, long *LDB);
/*~~~~~

void reg(long M, long N, double *QR, double *leadu, double *E,
        double *y, double *B, double *cov, double *se, double *e, double *sigma);

void read_name(char *);

main()
{
    char *filename;
    double *x, *y, *leadu, *E, *B, *cov, *se, *e, *sigma;
    long i, j, nrow, ncol, mindim;

    filename = calloc(MAX_NAME);

    /* matrix must be of the form [X,y] where first column of X is
    a vector of 1's if an intercept term is desired */

```

```

printf("\n%s\n%s", "Enter file name containing the matrix [X,y] ",
      "(with intercept column included if desired): ");
read_name(filename);
printf("\nEnter the number of observations: ");
scanf("%u",&nrow);
printf("\nEnter the number of parameters (including intercept): ");
scanf("%u",&ncol);
if(nrow <= ncol)
    printf("\n\nWARNING: #obs = &d <= &d = #parameters\n\n",nrow,ncol);
mindim = lmin(nrow,ncol); /* mindim is the smaller dimension */

x = dmalloc(nrow*(ncol+1));
y = x+(nrow*ncol); /* y is assigned the address of last col of x */
leadu = dmalloc(mindim);
E = dmalloc(ncol*ncol);
B = dmalloc(ncol);
cov = dmalloc(ncol*ncol);
se = dmalloc(ncol);
e = dmalloc(nrow);
sigma = dmalloc((long)1);

matread(x, nrow, ncol+1, filename);
/* matlabread(x, nrow, ncol, filename); */
/*matrix is stored contiguously column-wise */

qrpivot(nrow,ncol,x,E,leadu); /* only send first ncol columns of x */

printf("\nRough estimate of smallest singular value of X:");
printf("\nR(%d,%d) = %lf",ncol,ncol,x[nrow*(ncol-1)+(ncol-1)]);

reg(nrow,ncol,x,leadu,E,y,B,cov,se,e,sigma);

printf("\n\nMSE = %lf\n",*sigma);
printf("\nCOEFFICIENT \t SE \n");
for (i = 0; i < ncol; i++)
    printf("%4.5lf \t %4.5lf\n", B[i],se[i]);
printf("\n\nOBS \t RESIDUAL\n");
for (i = 0; i < nrow; i++)
    printf("%d \t %4.5lf\n", i+1,e[i]);
}

/* Subroutine reg()
Arguments:

    M number of rows of X

    N number of columns of X (expect N < M)

```

```

QR the matrix resulting from applying qrpivot() to X

leadu
    on entry: the vector of leading u's resulting from qrpivot()
    on exit: the vector of coefficient estimates B, where  $y = XB$ 

E the permutation matrix resulting from qrpivot()

y  a vector (length M) of "observables" (the rhs in  $XB = y$ )

B    on entry: an arbitrary length N vector
    on exit: the coefficient estimates

cov  on entry: an arbitrary NxN matrix
    on exit: the covariance matrix

se   on entry: an arbitrary length N vector
    on exit: the s.e.'s of the coefficient estimates

e    on entry: an arbitrary length M vector
    on exit: the vector of residuals:  $e = y - XB$ 

sigma on exit: the mse =  $y^te / (M-N)$ 
    */

void reg(long M, long N, double *QR, double *leadu, double *E,
        double *y, double *B, double *cov, double *se, double *e, double *sigma)
{
    long i,j,mindim;
    double a, *Qy, *invR, *EiR, *coef;

    /* BLAS arguments */
    long INC=(long)1;
    double alpha = (double)1, beta = (double)0;
    char UPLO, NOTRANS, TRANS, DIAG, SIDE;
    UPLO='U'; NOTRANS = 'N'; TRANS = 'T'; DIAG = 'N'; SIDE='L';

    dcopy_(&M, y, &INC, e, &INC);    /* e <- y */

    mindim = lmin(M-1,N);    /* expect mindim = N */

    /* Apply P(n)...P(1) to e to get e <- (Q_1 Q_2)^t Y */
    for(j=0;j<mindim;j++)
    {
        a = leadu[j]*e[j];    /* initialize a = u(1)e(1) */
        for(i=j+1;i<M;i++)
            a += QR[M*j+i]*e[i]; /* a = u^t e */
    }

```

```

        a *= (double)(-2);
        e[j] += a * leadu[j];      /* e(1) <- e(1) - 2 u(1)ute */
        for(i=j+1;i<M;i++)
            e[i] += a * QR[M*j+i]; /* e <- e + (-2)uut e */
    }

/*~~~~~
COMPUTE COEFFICIENTS

BLAS 3 method:      (currently used method) */

/* compute inv(R) */
invR = dmalloc(N*N);      /* workspace */
for(j=0;j<N;j++)          /* begin with identity matrix */
{
    for(i=0;i<N;i++)
        invR[N*j+i]=(double)0;
    invR[N*j+j]=(double)1;
}

/* invR <- alpha*inv(R)*invR = alpha*inv(R)*eye */
dtrsm_(&SIDE, &UPLO, &NOTRANS, &DIAG, &N, &N, &alpha, QR, &M,invR,&N);

/* compute the E*inv(R) matrix */
EiR = dmalloc(N*N);
/* EiR <- (alpha)E*invR + (beta)EiR */
dgemm_(&NOTRANS, &NOTRANS, &N, &N, &N, &alpha, E, &N,
      invR, &N, &beta, EiR, &N);
free(invR);

/* compute the coefficients */
dgemv_(&NOTRANS, &N, &N, &alpha, EiR, &N, e, &INC, &beta, B, &INC);
/* B <- (alpha)EiR*e + (beta)B      (beta = 0) only references 1st N elements of e */

/*~~~~~*/

/*~~~~~
COMPUTE COEFFICIENTS

Alternative method (BLAS 2):      */

if(0) /* not currently used */
{
    coef = dmalloc(N); /* workspace */
    dcopy_(&N, e, &INC, coef, &INC); /* coef <- e(1:N) */
    dtrsv_(&UPLO, &NOTRANS, &DIAG, &N, QR, &M, coef, &INC); /* coef <- Inv(R)*coef */
    dgemv_(&NOTRANS, &N, &N, &alpha, E, &N, coef, &INC, &beta, B, &INC);
    /* B <- (alpha)E*coef + (beta)B      (beta = 0) */
    free(coef);
}

```

```

    }

/*~~~~~*/

/*~~~~~
COMPUTE RESIDUALS
*/
    for(i=0;i<N;i++) e[i]=(double)0; /* annihilate first N elements of e */

    /* Apply P(1)...P(n) to e to get e <- Q2 Q2^t Y*/
    for(j=(mindim-1);j>=0;j--)
    {
        a = leadu[j]*e[j];          /* initialize a = u(1)e(1) */
        for(i=j+1;i<M;i++)
            a += QR[M*j+i]*e[i]; /* a = u^t e */
        a *= (double)(-2);
        e[j] += a * leadu[j];      /* e(1) <- e(1) - 2 u(1)u^te */
        for(i=j+1;i<M;i++)
            e[i] += a * QR[M*j+i]; /* e <- e + (-2)uu^t e */
    }
/*~~~~~*/

/*~~~~~
COMPUTE MSE
*/
    *sigma = ddot_(&M, y, &INC, e, &INC);
    *sigma /= (M - N);

/*~~~~~
COMPUTE COVARIANCE MATRIX and SE's

cov <- (alpha)EiR*(EiR)' + (beta)cov */
    dgemm_(&NOTRANS, &TRANS, &N, &N, &N, &alpha, EiR, &N,
        EiR, &N, &beta, cov, &N);

    for(j=0;j<N;j++)
        se[j] = sqrt((*sigma)*cov[N*j+j]);
}

void read_name(char *name)
{
    int c, i = 0;

    while ((c = getchar()) != EOF && c != ' ' && c != '\n')
        name[i++] = c;
    name[i] = '\0';
}

```

References

- [1] J. Demmel, *Applied Numerical Linear Algebra*. SIAM, Philadelphia, 1997.
- [2] A. Kelley and I. Pohl, *C By Dissection*. Addison Wesley, Menlo Park, 1996.
- [3] M. Loukides and A. Oram *Programming with GNU Software*. O'Reilly and Associates, Inc., Sebastopol, 1997.