

# Statistics 243: *assignment 4*

William J. De Meo

August 8, 2011

## 1 Sweep

The program `sweep.c` contains the subroutine `sweep()`, which performs the sweep algorithm on a symmetric matrix (actually, the function `colsweep` does all the work).<sup>1</sup> The program is written for a symmetric matrix because that is what we find in applications like regression, and because foresight that the matrix is symmetric affords tremendous savings in storage and computational costs. The `sweep()` routine accesses only the upper triangle of its matrix argument, and overwrites only this portion of the matrix. Therefore, if an entire symmetric matrix is passed to `sweep()`, then on exit the same matrix has the answer in its upper triangle and the original matrix in the lower triangle. The diagonal entries are the negative reciprocals of the original diagonal entries.

Using MATLAB, we show in section 3.1 that, when sweeping all columns of a matrix, we get the negative of the inverse (or close to it). When sweeping the first five columns of the 6 by 6 matrix  $X^t X$ , where  $X = [X : y]$  is an augmented matrix of covariates and response variable, we arrive at the matrix:

$$\begin{pmatrix} -(X^t X)^{-1} & \hat{\beta} \\ \hat{\beta}^t & y^t(I - P)y \end{pmatrix}$$

where  $\hat{\beta} = (X^t X)^{-1} X^t y$  and  $P \equiv X(X^t X)^{-1} X^t$  is the projection operator.

## 2 Non-linear Regression

### 2.1 Parameter Estimates

Write the non-linear regression model as:<sup>2</sup>

$$y_n = f(x_n, \theta) + z_n$$

where  $f$  is the expectation function and  $x_n$  is a vector of associated covariates for the  $n$ th observation. Further consider the  $N$  row vectors  $x_n$ , ( $n = 1, \dots, N$ ) as fixed and concentrate on the relation between the response  $y_n$  and the parameters  $\theta$ . Thus, letting  $\eta_n(\theta) = f(x_n, \theta)$ , we have

$$Y = \eta(\theta) + Z$$

and we assume

$$E[Z] = 0, \quad \text{Var}(Z) = E[ZZ^t] = \sigma^2 I$$

---

<sup>1</sup>The algorithm used is similar to that found in [2], p. 87

<sup>2</sup>For further details, consult Bates and Watts [1].

as in the linear model. We refer to  $\eta(\theta)$  as the *expectation surface*. When  $\theta \in \Theta \subset \mathbf{R}^p$ ,  $\eta$  is generally a  $p$  dimensional curved surface in the  $N$  dimensional covariate space.

Having set up the context, we seek the least squares estimate of  $\theta$ . That is, we want  $\hat{\theta} \in \Theta$  such that,

$$\|Y - \eta(\hat{\theta})\|_2 = \min_{\theta \in \Theta} \|Y - \eta(\theta)\|_2$$

This estimate is found by solving a system of non-linear equations, which requires the use of iterative procedures. The subroutine `gn()` applied in section 3.2 employs the familiar Gauss-Newton procedure, which we describe briefly as follows: start with an initial guess  $\theta_0$ , and write the Taylor approximation (ignoring small order terms) as

$$\eta(\theta) \approx \eta(\theta_0) + J_0(\theta - \theta_0)$$

Now note that this is equivalent to approximating the residuals  $z(\theta) = Y - \eta(\theta)$  by

$$z(\theta) \approx (Y - \eta(\theta_0)) - J_0(\theta - \theta_0) \equiv z_0 - J_0\delta$$

Since  $\theta_0$  is our known starting value, estimating  $\theta - \theta_0$  is equivalent to estimating  $\theta$ . Thus we seek that  $\delta_0 \equiv (\theta - \theta_0)$  which minimizes the residual:

$$\|z(\theta)\|_2 \equiv \|z_0 - J_0\delta_0\|_2 \tag{1}$$

and take as our new estimate  $\theta_1 = \theta_0 + \hat{\delta}_0$ . Proceeding inductively, the typical step of the algorithm is:

$$\begin{aligned} \theta_{i+1} &= \theta_i + (J_i^t J_i)^{-1} J_i^t [Y - \eta(\theta_i)] \\ &= \theta_i + \hat{\delta}_i \end{aligned}$$

We have reduced each iteration to the usual least squares problem, to which we could apply the `sweep()` routine from the problem above. However, the QR routine submitted for assignment three has a number of advantages; perhaps most importantly, it warns us when  $J_i$  is nearly rank deficient. Therefore, we use the QR routine to compute the successive  $\theta_{i+1} = \theta_i + \hat{\delta}_i$ .

## 2.2 A Measure of Uncertainty

It has been said that an estimate unaccompanied by a measure of uncertainty is almost worthless. Thus, the includes estimated standard errors for the  $\hat{\theta}$ . Since exact expressions are unavailable in the non-linear case, measures of uncertainty for the Gauss-Newton algorithm are based on an *asymptotic* covariance matrix, which is estimated by:

$$\widehat{Cov}(\hat{\theta}) = s^2 (J^t J)^{-1}$$

where  $J$  is the Jacobian matrix from the discussion above, evaluated at the least squares estimates obtained upon the final iteration, and  $s^2$  is the residual mean square given by

$$s^2 = \|Y - \eta(\hat{\theta})\|^2 / (N - p)$$

In the example analyzed in section 3.2,  $v$  is the first parameter, so the approximate standard error of  $\hat{v}$  is the square root of the (1,1) entry of  $\widehat{Cov}(\hat{\theta})$ . That is

$$s_v = \sqrt{s^2 e_1^t (J^t J)^{-1} e_1}$$

where  $e_1 = (1, 0, \dots, 0)^t$ . These are output by the program `nlstest.c` in section 3.2.

## 3 Results

### 3.1 Sweep

In the following interaction with the `sweepTest.c`, which calls the `sweep()` routine, there are two features to notice: 1) the user can specify which, and how many of the columns to sweep, and 2) it doesn't matter what appears in the lower triangle of the matrix (we get the same answers using the full SSCP matrix, as we do when using SSCPU).

The matrix SSCP (called `S` below) was produced in MATLAB. It is the matrix  $X^t X$ , where  $X$  is a  $100 \times 6$  matrix of normal (0,1) random numbers.

```
% sweepTest
```

```
Enter name of the file containing the spd matrix: SSCP
```

```
Enter the dimension: 6
```

```
Enter number of columns to be swept: 6
```

```
column for sweep 0: 0
```

```
column for sweep 1: 1
```

```
column for sweep 2: 2
```

```
column for sweep 3: 3
```

```
column for sweep 4: 4
```

```
column for sweep 5: 5
```

```
The matrix is:
```

74.90517	-2.05071	0.68651	7.51039	-5.73764	0.32840
-2.05071	89.96492	1.30913	3.03349	-12.17186	-1.65234
0.68651	1.30913	91.26154	-0.32621	-18.66620	19.21324
7.51039	3.03349	-0.32621	98.56231	2.74726	-10.59074
-5.73764	-12.17186	-18.66620	2.74726	94.53465	-13.85515
0.32840	-1.65234	19.21324	-10.59074	-13.85515	95.10084

```
The sweep produced:
```

-0.01354	-0.00047	-0.00008	0.00107	-0.00093	-0.00000
-2.05071	-0.01136	-0.00007	0.00039	-0.00159	-0.00041
0.68651	1.30913	-0.01176	0.00023	-0.00199	0.00188
7.51039	3.03349	-0.32621	-0.01036	0.00026	-0.00108
-5.73764	-12.17186	-18.66620	2.74726	-0.01155	-0.00165
0.32840	-1.65234	19.21324	-10.59074	-13.85515	-0.01052

Now we try it with only the upper triangle:

```
% sweeptest
```

```
Enter name of the file containing the spd matrix: SSCPU
```

```
Enter the dimension: 6
```

```
Enter number of columns to be swept: 6
```

```
column for sweep 0: 0
```

```
column for sweep 1: 1
```

```
column for sweep 2: 2
```

```
column for sweep 3: 3
```

```
column for sweep 4: 4
```

```
column for sweep 5: 5
```

```
The matrix is:
```

74.90517	-2.05071	0.68651	7.51039	-5.73764	0.32840
0.00000	89.96492	1.30913	3.03349	-12.17186	-1.65234
0.00000	0.00000	91.26154	-0.32621	-18.66620	19.21324
0.00000	0.00000	0.00000	98.56231	2.74726	-10.59074
0.00000	0.00000	0.00000	0.00000	94.53465	-13.85515
0.00000	0.00000	0.00000	0.00000	0.00000	95.10084

```
The sweep produced:
```

-0.01354	-0.00047	-0.00008	0.00107	-0.00093	-0.00000
0.00000	-0.01136	-0.00007	0.00039	-0.00159	-0.00041
0.00000	0.00000	-0.01176	0.00023	-0.00199	0.00188
0.00000	0.00000	0.00000	-0.01036	0.00026	-0.00108
0.00000	0.00000	0.00000	0.00000	-0.01155	-0.00165
0.00000	0.00000	0.00000	0.00000	0.00000	-0.01052

To verify that, after sweeping all 6 columns, we have the negative of the inverse, we take the inverse using MATLAB:

```
>> S
```

```
S =
```

74.9052	-2.0507	0.6865	7.5104	-5.7376	0.3284
-2.0507	89.9649	1.3091	3.0335	-12.1719	-1.6523
0.6865	1.3091	91.2615	-0.3262	-18.6662	19.2132
7.5104	3.0335	-0.3262	98.5623	2.7473	-10.5907
-5.7376	-12.1719	-18.6662	2.7473	94.5347	-13.8551

```

0.3284   -1.6523   19.2132  -10.5907  -13.8551   95.1008

>> format long

>> inv(S)

ans =

Columns 1 through 4

    0.01354111561666    0.00046799568625    0.00008525124555   -0.00107611827999
    0.00046799568625    0.01135736256171    0.00007644175507   -0.00038944059053
    0.00008525124555    0.00007644175507    0.01181901254893   -0.00025423412626
   -0.00107611827999   -0.00038944059053   -0.00025423412626    0.01037194437930
    0.00092420563087    0.00157074269915    0.00204583692967   -0.00029704847656
   -0.00004104512204    0.00036574088772   -0.00211701973720    0.00116008917857

Columns 5 through 6

    0.00092420563087   -0.00004104512204
    0.00157074269915    0.00036574088772
    0.00204583692967   -0.00211701973720
   -0.00029704847656    0.00116008917857
    0.01143124849618    0.00124310603658
    0.00124310603658    0.01125965024054

```

The results are nearly the same. We suspect the difference is due to round-off error, and not the instability of our method of taking inverses. The condition number of  $S$  is  $\|S\|_2\|S^{-1}\|_2 \approx 1.98$ ,

```
>> cond(S)
```

```
ans =
```

```
1.97781040735287
```

Let's try sweeping only the first 5 columns:

```
% sweeptest
```

```
Enter name of the file containing the spd matrix: SSCP
```

```
Enter the dimension: 6
```

```
Enter number of columns to be swept: 5
```

```
column for sweep 0: 0
```

```
column for sweep 1: 1
```

```
column for sweep 2: 2
```

column for sweep 3: 3

column for sweep 4: 4

The matrix is:

74.90517	-2.05071	0.68651	7.51039	-5.73764	0.32840
-2.05071	89.96492	1.30913	3.03349	-12.17186	-1.65234
0.68651	1.30913	91.26154	-0.32621	-18.66620	19.21324
7.51039	3.03349	-0.32621	98.56231	2.74726	-10.59074
-5.73764	-12.17186	-18.66620	2.74726	94.53465	-13.85515
0.32840	-1.65234	19.21324	-10.59074	-13.85515	95.10084

The sweep produced:

-0.01335	0.00006	-0.00002	-0.00033	0.00023	0.00438
-2.05071	-0.01112	-0.00000	-0.00003	0.00007	-0.01837
0.68651	1.30913	-0.01096	0.00000	0.00008	0.21053
7.51039	3.03349	-0.32621	-0.01015	0.00001	-0.10745
-5.73764	-12.17186	-18.66620	2.74726	-0.01058	-0.14656
0.32840	-1.65234	19.21324	-10.59074	-13.85515	95.10084

That is, our OLS estimate is  $\hat{\beta} = (0.00438, -0.01837, 0.21053, -0.10745, -0.14656)$ , as compared to that given by MATLAB:

```
>> Y = X(:,p)
```

```
>> X = X(:,1:5)
```

```
>> [beta,se]=lskov(X,Y,eye(100))
```

beta =

```
0.00364532833282
-0.03248243772267
0.18801825029833
-0.10303065848312
-0.11040361023850
```

se =

```
0.11251244721632
0.10298814376501
0.10333021250112
0.09790140713222
0.10275423659730
```

My estimates are slightly different and, at this point, I can't say why. Notice also that the se given by MATLAB is readily computable from our sweep output. MATLAB computes it using  $se = \sqrt{\text{diag}(X^t X)^{-1} \text{mse}}$ , and each component of this expression appears in the result of our sweep program.

### 3.2 Non-linear Regression

A description of the first problem on which we test the gn.c program was given as follows:<sup>3</sup>

An experiment was carried out wherein varying doses of Rose Bengal, an enzyme, was used as a catalyst for the production of an protein substrate. The concentration of the substrate was measured spectrophotometrically. The first column of the data represents the dose of the enzyme, and the second column the spectrophotometric response. The equation proposed for the data is

$$\text{response} = \frac{\theta_1 \cdot \text{dose}}{\theta_2 + \text{dose}} \quad (2)$$

where  $\theta_1$  and  $\theta_2$  are the parameters to be estimated. (This equation is known as the Michaelis-Menten model for the velocity of an enzyme catalyzed reaction.)

Such non-linear functions are rather basic and easy to change, so we hard code it into the calling program `nlstest.c`. Now we need to decide on good initial values for  $\theta = (\theta_1, \theta_2)$ . According to the set up presented in section 2, the function given by (2) corresponds to the expectation surface,  $\eta(\theta)$ . That is,

$$E_\theta[Y] = \eta(\theta) = \frac{\theta_1 \cdot \text{dose}}{\theta_2 + \text{dose}}$$

We estimate  $E_\theta[Y]$  from the data by  $\bar{Y} = \sum y_i / N = 23.4$ . Looking at a scatter plot of the data, we notice approximately, as  $\text{dose} \rightarrow \infty$ ,  $\text{response} \rightarrow 30$ . Therefore, L'Hopital's rule would suggest an initial estimate of  $\theta_1 = 30$ . For a crude estimate of  $\theta_2$ , we compute the sample mean  $\sum \text{dose} / N \equiv \bar{d} = 0.228875$  and equate

$$\bar{Y} = \frac{\theta_1 \bar{d}}{\theta_2 + \bar{d}}$$

yielding the estimate  $\theta_2 = 0.065$ . Therefore, we take as our initial guess,  $\theta_0 = (30, 0.065)$ . Invoking the program `nlstest` with these initial values produces the following:

```
% nlstest

Enter file name containing the data: prob1.dat

gn() exited with return value 1, on iteration 7

MSE = 0.214164

COEFFICIENT      SE
33.12465         0.37729
0.04606          0.00208
```

---

<sup>3</sup>Source: BMDP manual, p. 395.

## 4 Program Listings

### 4.1 sweep program listings

```

/*****
 * sweeptest.c  main program for testing sweep subroutine  *
 *                                                    *
 * Created by William J. De Meo                        *
 * on 12/17/97                                          *
 *                                                    *
 *****/
#include <stdlib.h>
#include <stdio.h>
#include "prototypes.h"
#define MAX_NAME 100

void read_name(char *);

main()
{
    char *filename;
    double *x;
    long i, k, nrow, *index;

    filename = calloc(MAX_NAME);

    printf("\nEnter name of the file containing the symmetric matrix: ");
    read_name(filename);
    printf("\nEnter the dimension: ");
    scanf("%u",&nrow);

    printf("\nEnter number of columns to be swept: ");
    scanf("%u",&k);
    index = malloc(k);

    for(i=0;i<k;i++)
    {
        printf("\ncolumn for sweep %d: ",i);
        scanf("%u",index+i);
    }

    x = malloc(nrow*nrow);
    /* matread(x, nrow, ncol, filename); */
    matlabread(x, nrow, nrow, filename);
    /*matrix is stored contiguously column-wise */
    printf("\nThe matrix is: \n");
    matprint(x,nrow,nrow);

```



```

    /* Test sweep: */
    sweep(nrow,x,k,index);

    printf("\nThe sweep produced: \n");
    matprint(x,nrow,nrow);
}

void read_name(char *name)
{
    int c, i = 0;

    while ((c = getchar()) != EOF && c != ' ' && c != '\n')
        name[i++] = c;
    name[i] = '\0';
}

/*~~~~~
sweep.c

Created on 12/16/97 by William J. De Meo
Last modified 12/17/97

Purpose: Perform sweep algorithm of an M dimensional symmetric matrix

Further Details: This implementation uses BLAS 1
                  Requires subroutines found in the libraries:
                  sunperf, and blas
                  the later two are linked with the options:
                  -lsunperf -dalign -lblas
                  compilation must be done with the -dalign option
~~~~~*/
#include "prototypes.h"

void daxpy_(long *N, double *alpha, double *x, long *INCX, double *Y, long *INCY);
/* Y <- alpha X + Y */

/* Subroutine sweep: perform a sweep of p columns of A

Arguments:
    M (long) dimension of A

    A (pointer to double)
    on entry: the M by M symmetric matrix to be swept
              (lower triangle is never accessed)
    on exit: upper triangle = the swept matrix
              (lower triangle is unchanged)

```

```

        p  (long) number of columns to be swept

        index  (pointer to long) a vector of length p containing the
                indices specifying which columns are to be swept
                Note: indices can take values in [0, M-1]
*/

void sweep(long M, double *A, long p, long *index)
{
    long i;
    long k;

    for(i=0;i<p;i++)
    {
        k = index[i];
        colsweep(M,A,k);
    }
}

/* Subroutine colsweep: sweep column k of A

Arguments:
    M  (long) dimension of A

    A  (pointer to double)
        on entry: the M by M symmetric matrix to be swept
                 (lower triangle is never accessed)
        on exit: upper triangle = the swept matrix
                 (lower triangle is unchanged)

    k  (long) column to be swept
*/

void colsweep(long M, double *A, long k)
{
    long i,j,length,one = (long)1;
    double alpha;

    /* First convert non-k elements */

    /* rows 0 to k-1 */
    for(i=0;i<k;i++)
    {
        length = k-i;

```

```

/* i < k, i <= j < k (triangle) */
/* for columns i:k of rows i:k, do the following daxpy: */
/* a(i,i:M) <- alpha*a(k,i:M) + a(i,i:M) */
alpha = ((double)-1)*A[k*M+i]/A[k*M+k];
daxpy_(&length,&alpha,A+(k*M+i),&one,A+(i*M+i),&M);
/* notice: replacing row k col i here ^ with row i col k since sym */

length = M-k-1;
/* i < k, j > k (rectangle) */
daxpy_(&length,&alpha,A+((k+1)*M+k),&M,A+((k+1)*M+i),&M);

}

/* rows k+1 to M */
for(i=k+1;i<M;i++)
{
    length = M-i-1;

    /* k < i < j (triangle) */
    /* for columns k:M of rows k:M, do the following daxpy: */
    /* a(i,i:M) <- alpha*a(k,i:M) + a(i,i:M) */
    alpha = ((double)-1)*A[i*M+k]/A[k*M+k];
    daxpy_(&length,&alpha,A+(i*M+k),&M,A+(i*M+i),&M);
}

/* divide elements 0:k-1 of kth col of A by a_{kk} */
for(i=0;i<k;i++)    A[k*M+i]/=A[k*M+k];
/* divide elements k+1:M of the kth row of A by a_{kk} */
for(j=k+1;j<M;j++)    A[j*M+k]/=A[k*M+k];
/* Finally, invert diagonal */
A[k*M+k] = ((double)-1)/A[k*M+k];
}

```

## 4.2 non-linear regression program listings

\*\*\*\*\*

nlstest.c main program for testing gn() subroutine

Created by William J. De Meo on 12/20/97

last modified on 12/21/97

Purpose: Perform non-linear least squares using the gn()  
 implementation of the Gauss Newton algorithm found in  
 /accounts/grad/chip/lib/gauss/gn.c

The following describes the problem used for the test:

(Source: BMDP manual, p. 395)

An experiment was carried out wherein varying doses of Rose Bengal, an enzyme, was used as a catalyst for the production of an protein substrate. The concentration of the substrate was measured spectrophotometrically. The first column of the data represents the dose of the enzyme, and the second column the spectrophotometric response. The equation proposed for the data is

$$\text{response} = \frac{v * \text{dose}}{(k + \text{dose})},$$

where  $v$  and  $k$  are the parameters to estimated. (This equation is known as the Michaelis-Menten model for the velocity of an enzyme catalyzed reaction.)

The data are:

dose	response
0.027	12.7
0.044	16.0
0.073	20.4
0.102	22.3
0.175	26.0
0.257	28.8
0.483	29.6
0.670	31.4

```

*****/
#include <stdlib.h>
#include <stdio.h>
#include <math.h>
#include "prototypes.h"
#define MAX_NAME 100

/* Data specific settings: */
/* number of observations */
#define M 8
/* number of covariates */
#define N 1
/* number of parameters */
#define P 2

#define MAXITER 20

void init(double *);
void eta(long, long, double *, long, double *, double *);

```

```

void jac(long, long, double *, long , double *, double *);
void (*f)(long, long, double *, long , double *, double *);
void (*g)(long, long, double *, long , double *, double *);

void gn(long , long, double *X, double *Y, long, double *theta,
        void (*f)(long , long , double *X, long , double *theta, double *F),
        void (*g)(long , long , double *X, long , double *theta, double *J),
        double *cov, double *sigma, long iter);

void read_name(char *name);

main()
{
    double *x, *y, *theta, *cov, *sigma, *se;
    long i, j;
    char *filename;

    filename = calloc(MAX_NAME);
    printf("\nEnter file name containing the data: ");
    read_name(filename);

    x = malloc(M*(N+1));
    y = x+M; /* y is assigned the address of second col of x */
    theta = malloc(P);
    se = malloc(P);
    cov = malloc(P*P);
    sigma = malloc((long)1);

    matread(x, M, N+1, filename);

    init(theta);

    f = eta;
    g = jac;

    gn(M,N,x,y,P,theta,f,g,cov,sigma,MAXITER);

    /* compute se's: */
    for(j=0;j<P;j++)
        se[j] = sqrt((*sigma)*cov[P*j+j]);
    printf("\n\nMSE = %lf\n",*sigma);
    printf("\nCOEFFICIENT \t SE \n");
    for (i = 0; i < P; i++)
        printf("%4.5lf \t %4.5lf\n", theta[i],se[i]);
}

void init(double *theta)
{

```

```

    theta[0]=(double)30;
    theta[1]=(double)0.065;
}

void eta(long m, long n, double *x, long p, double *th, double *F)
{
    long i;
    for(i=0;i<m;i++)
        F[i] = (th[0] * x[i])/(th[1] + x[i]);
}

void jac(long m, long n, double *x, long p, double *th, double *J)
{
    long i;

    for(i=0;i<m;i++)
    {
        J[i] = x[i]/(th[1] + x[i]);
        J[M+i] = -th[0]*x[i]/pow((th[1] + x[i]),2);
    }
}

void read_name(char *name)
{
    int c, i = 0;

    while ((c = getchar()) != EOF && c != ' ' && c != '\n')
        name[i++] = c;
    name[i] = '\0';
}

/*~~~~~
gn.c

Created on 12/18/97 by William J. De Meo
Last modified 12/21/97

Purpose: Perform one iteration of Gauss-Newton with step halving

Further Details: This implementation uses BLAS 2
                 (matrix-vector mult. and rank 1 updates)

Dependencies:
subroutines found in:
    /accounts/grad/chip/lib/decomp/QR/House.c
    sunperf, and blas
the later two are linked with the options:

```

```

        -lsunperf -dalign -lblas
        compilation must be done with the -dalign option
~~~~~*/

#include <math.h>

/* ~~~Prototypes ~~~~~

For subroutines in /accounts/grad/chip/lib/decomp/QR/House.c:      */
#include "prototypes.h"

/*
For BLAS subroutines:      */

double dnorm2(long *N, double *x, long *INC); /* L2 norm of x */
void dcopy_(long *N, double *X, long *INCX, double *Y, long *INCY); /* y <- x */

/* y <- (alpha)Ax + (beta)y (or A^t if TRANSA='T') */
void dgemv_(char *TRANSA, long *M, long *N, double *alpha, double *A, long *LDA,
            double *x, long *INCX, double *beta, double *y, long *INCY);

/* C <- (alpha)AB + (beta)C */
void dgemm_(char *TRANSA, char *TRANSB, long *M, long *N, long *K, double *alpha,
            double *A, long *LDA, double *B, long *LDB, double *beta, double *C, long *LDC);

/* B <- alpha*inv(A)*B */
void dtrsm_(char *SIDE, char *UPLO, char *TRANSA, char *DIAG, long *M, long *N,
            double *alpha, double *A, long *LDA, double *B, long *LDB);
/* ~~~~~*/

void reg(long M, long N, double *QR, double *leadu, double *E,
        double *y, double *B, double *cov, double *se, double *e, double *sigma);

void read_name(char *);

void Exit(int i, long k);
void Free(double*, double*, double*, double*, double*, double*, double*, double*);

/* Tolerance for smallest singular value of J */
#define SINGTOL 1.0e-7
#define DELTOL 1.0e-7
#define STEPTOL 1.0e-4

/* ~~~~~
Subroutine gn:

Arguments:
    M (long) number of rows of X

```

```

N (long) number of columns of X

X (pointer to double) the M by N matrix of covariates

Y (pointer to double) M by 1 observable vector

P (long) number of parameters

theta (pointer to double) P-vector of parameters
    on entry: initial guess
    on exit: improved solution after one iteration

(*f)(M,N,X,P,theta,F) points to a function which evaluates to the expectation
    surface for given covariates X and parameter values theta

(*g)(M,N,X,P,theta,J) points to a function which evaluates to the analytic
    Jacobian of f

cov    on entry: an arbitrary PxP matrix
    on exit: the estimated asymptotic covariance matrix

sigma  on exit: the mse = |y - f(theta)|^2 / (M-P)

iter (long) max number of iterations

*/
void gn(long M, long N, double *X, double *Y, long P, double *theta,
    void (*f)(long M, long N, double *X, long P, double *theta, double *F),
    void (*g)(long M, long N, double *X, long P, double *theta, double *J),
    double *cov, double *sigma, long iter)
{
    long i,j,k, one=(long)1;
    double *F, *J, *leadu, *E, *z, *delta, newS, oldS, half;
    double *pwork,*ework,test;

    F = dmalloc(M);
    J = dmalloc(M*P);
    z = dmalloc(M);
    E = dmalloc(P*P);
    leadu = dmalloc(P);
    delta = dmalloc(P);

    pwork = dmalloc(P); /* parameter workspace */
    ework = dmalloc(M); /* error workspace */

    for(k=0;k<iter;k++)

```



```

{
    (*f)(M,N,X,P,theta,F);
    /* z <- y - f(theta) */
    for(i=0;i<M;i++) z[i] = Y[i] - F[i];

    oldS = dnorm2_(&M,z,&one); /* L2 norm of error */

    /* Form Jacobian and QR decomposition */
    (*g)(M,N,X,P,theta,J);

    qrpivot(M,P,J,E,leadu);
    /*printf("\nJ[%d] = \n",k+1);
    matprint(J,M,P);*/

    if(fabs(J[M*(P-1)+(P-1)]) < SINGTOL)
    {
        Exit(0,k); /* 0 <= J rank defficient (bad exit) */
        Free(F,J,z,E,leadu,delta,pwork,ework);
        return;
    }

    /* reg for fitting z ~ J.delta */
    reg(M,P,J,leadu,E,z,delta,cov,pwork,ework,sigma);

    /* Don't need errors=pwork or se=ework computed by reg() */
    /* So P vector pwork and M vector ework are used
    for a different purpose below */

    /* cov = inv(J^tJ) for J corresponding to current theta
    so if we exit with codes 1 or 2 below, i.e. before computing a new theta,
    we return the correct covariance matrix */

    test=(double)0;
    for(j=0;j<P;j++)
    {
        if(fabs(delta[j])>test) test = delta[j];
    }
    if(test < DELTOL)
    {
        *sigma = oldS/(M-P); /* get new sigma before exit */
        Exit(1,k); /* 1 <= no change in theta_i */
        Free(F,J,z,E,leadu,delta,pwork,ework);
        return;
    }

    j=0;
    do
    {

```

```

        if((half = pow(.5,j)) < STEPTOL)
        {
            *sigma = oldS/(M-P);          /* get new sigma before exit */
            Exit(2,k); /* 2 <= No improvements from steps */
            Free(F,J,z,E,leadu,delta,pwork,ework);
            return;
        }
        j++;
        /* compute new candidate coefficients */
        for(i=0;i<P;i++) pwork[i] = theta[i] + half*delta[i];
        (*f)(M,N,X,P,pwork,F);
        for(i=0;i<M;i++) ework[i] = Y[i] - F[i];          /* compute new error */
        newS = dnorm2_(&M, ework, &one); /* L2 norm of new error*/
    }while(newS > oldS);

    dcopy_(&P,pwork, &one, theta, &one); /* theta <- pwork */
}

}/*~~~ end gn() ~~~~~*/

/*~~~~~
Subroutine reg()
Arguments:

    M number of rows of X

    N number of columns of X (expect N < M)

    QR the result of applying qrpivot() to X
        (i.e. upper triangle is R, lower trapezoid are u's)

    leadu
        on entry: the vector of leading u's resulting from qrpivot()
        on exit: the vector of coefficient estimates B, where y = XB

    E the permutation matrix resulting from qrpivot()

    y a vector (length M) of "observables" (the rhs in XB = y)

    B on entry: an arbitrary length N vector
        on exit: the coefficient estimates

    cov on entry: an arbitrary NxN matrix
        on exit: the covariance matrix

    se on entry: an arbitrary length N vector
        on exit: the s.e.'s of the coefficient estimates

```

```

    e      on entry: an arbitrary length M vector
           on exit: the vector of residuals:  e = y - XB

    sigma  on exit: the mse =  y^te / (M-N)
    */

void reg(long M, long N, double *QR, double *leadu, double *E,
        double *y, double *B, double *cov, double *se, double *e, double *sigma)
{
    long i,j,mindim;
    double a, *Qy, *invR, *EiR, *coef;

    /* BLAS arguments */
    long INC=(long)1;
    double one = (double)1, zero = (double)0;
    char UPLO, NOTRANS, TRANS, DIAG, SIDE;
    UPLO='U'; NOTRANS = 'N'; TRANS = 'T'; DIAG = 'N'; SIDE='L';

    dcopy_(&M, y, &INC, e, &INC);      /* e <- y */

    mindim = lmin(M-1,N);      /* expect mindim = N */

    /* Apply P(n)...P(1) to e to get e <- (Q_1 Q_2)^t Y*/
    for(j=0;j<mindim;j++)
    {
        a = leadu[j]*e[j];          /* initialize  a = u(1)e(1) */
        for(i=j+1;i<M;i++)
            a += QR[M*j+i]*e[i]; /* a = u^t e */
        a *= (double)(-2);
        e[j] += a * leadu[j];      /* e(1) <- e(1) - 2 u(1)u^te */
        for(i=j+1;i<M;i++)
            e[i] += a * QR[M*j+i]; /* e <- e + (-2)uu^t e */
    }
    /*~~~~~
    COMPUTE COEFFICIENTS

    /* compute inv(R) */
    invR = dmalloc(N*N);      /* workspace */
    for(j=0;j<N;j++)          /* begin with identity matrix */
    {
        for(i=0;i<N;i++)
            invR[N*j+i]=(double)0;
        invR[N*j+j]=(double)1;
    }
    /* invR <- one*inv(R)*invR = one*inv(R)*eye */
    dtrsm_(&SIDE, &UPLO, &NOTRANS, &DIAG, &N, &N, &one, QR, &M,invR,&N);

    /* compute the E*inv(R) matrix */

```

```

EiR = dmalloc(N*N);
/* EiR <- (one)E*invR + (zero)EiR */
dgemm_(&NOTRANS, &NOTRANS, &N, &N, &N, &one, E, &N,
      invR, &N, &zero, EiR, &N);
free(invR);

/* compute the coefficients */
dgemv_(&NOTRANS, &N, &N, &one, EiR, &N, e, &INC, &zero, B, &INC);
/* B <- (one)EiR*e + (zero)B      (zero = 0)  only references 1st N elements of e */

/***** residuals and mse for the OLS part not necessary *****/
if(0)
{
    /* COMPUTE RESIDUALS
    */
    for(i=0;i<N;i++) e[i]=(double)0; /* annihilate first N elements of e */

    /* Apply P(1)...P(n) to e  to get e <- Q2 Q2^t Y*/
    for(j=(mindim-1);j>=0;j--)
    {
        a = leadu[j]*e[j];          /* initialize  a = u(1)e(1) */
        for(i=j+1;i<M;i++)
            a += QR[M*j+i]*e[i]; /* a = u^t e */
        a *= (double)(-2);
        e[j] += a * leadu[j];      /* e(1) <- e(1) - 2 u(1)u^te */
        for(i=j+1;i<M;i++)
            e[i] += a * QR[M*j+i]; /* e <- e + (-2)uu^t e */
    }

    /* COMPUTE MSE  (wrong MSE in nonlinear case. currently embraced by if(0))
    */
    *sigma = ddot_(&M, y, &INC, e, &INC);
    *sigma /= (M - N);

}

/**** end if(0) ****/

/*~~~~~
COMPUTE COVARIANCE MATRIX and SE's

cov <- (one)EiR*(EiR)' + (zero)cov */
dgemm_(&NOTRANS, &TRANS, &N, &N, &N, &one, EiR, &N,
      EiR, &N, &zero, cov, &N);

/* wrong se's:
for(j=0;j<N;j++)
se[j] = sqrt((*sigma)*cov[N*j+j]);

```

```

    */
}/*~~~ end reg() ~~~~~*/

void Exit(int i, long k)
{
    printf("\n\ngn() exited with return value %d, on iteration %d",i,k+1);
}

void Free(double *a, double *b, double *c, double *d, double *e,
          double *f,double *g, double *h)
{
    free(a);free(b);free(c);free(d);free(e);free(f);free(g);free(h);
}

```

## References

- [1] D. Bates and D. Watts, *Nonlinear Regression Analysis and its Applications*. John Wiley, New York, 1988.
- [2] R. Thisted, *Elements of Statistical Computing*. Chapman and Hall, New York, 1988.